

SIDDIGAN TM

Data Engineer

✉ 2001siddi@gmail.com  Siddigan TM  siddigan ☎ 6382771052 📍 Gurgaon

Data Engineer with 2+ years of experience designing cloud-native data lakes on AWS and Snowflake, building scalable data pipelines, and delivering curated analytics datasets across pharma and financial domains. Skilled in modern data engineering and AI workflows.

PROFESSIONAL EXPERIENCE

Data Engineer, ZS

2025 – present
Gurgaon, India

- Designed and implemented an end-to-end **pharma data lake architecture** on **AWS and Snowflake**, delivering ingestion, curated data models, and analytics ready datasets for **10+ downstream analytics consumers**.
- **Led requirements gathering, data profiling, and foundational data modeling**, defining KPI metrics and building ARD/RRD reporting views used by **business and operations teams across multiple regions**.
- **Engineered and deployed scalable, end-to-end ETL pipelines** using Apache Airflow, AWS EMR, AWS Glue, PySpark, and SparkSQL, processing around **billion records/day** ensuring efficient ingestion & transformation.
- **Migrated a legacy data lake** into the cloud by rerouting pipelines and implementing co loading strategies, enabling **zero downtime transition** and reducing batch processing delays by **~25%**.
- **Enhanced and optimized existing data products with evolving pharmaceutical business needs** and proactively addressing potential future use cases such as clinical trial product launch and real-world evidence analytics.

Data Engineer - Apprentice, Fidelity Investments

2024 – 2025
Bengaluru, India

- **Played a key role in the cloud migration of enterprise data lake to the cloud**, designing and optimizing end-to-end data pipelines for the Data Lake. Transitioned data Oracle DB systems into Snowflake with AWS.
- **Developed and implemented curated tables, stored procedures, and dynamic scripts**, leveraged Control-M, AWS Lamda and Apache NiFi to streamline workflows and scale data processing across enterprise systems.
- **Created and maintained operational Power BI dashboards**, improved operational efficiency and delivered business impact that reduced failure rate by 78%.
- **Worked on a POC for data masking Snowflake DB using Tonic AI**, which included designing, building, testing, and conducting thorough risk and impact analyses to support data privacy and compliance initiatives for data governance.
- **Built robust ETL testing and reconciliation frameworks with ICEDQ**, including functional, regression, and data quality validations, ensuring accuracy, consistency, and compliance with pharma industry standards.

SKILLS

Programming: Python, SQL, PL/SQL

Cloud & Data Platforms: AWS (S3, EMR, Glue, Athena), Snowflake

Data Engineering & ETL: Apache Airflow, Spark (PySpark, SparkSQL), Apache Kafka

Data Core: Dimensional Modeling, ETL/ELT, Data Warehousing, Power BI & Data Analysis

DevOps & Tools: Docker, Git, Jenkins, CI/CD

EDUCATION

Bachelor of Technology: Computer Science Engineering, *Lovely Professional University*
Machine Learning Specialization

2019 – 2023
Phagwara

PROJECTS

Financial Advisory and Consultancy Chatbot • Developed a financial advisory chatbot using a fine-tuned Llama 3 model integrated with a Retrieval-Augmented Generation (RAG) system, fetching and analysing real-time stock metrics and market news for personalized investment advice. • Engineered data pipelines using Apache NiFi and Airflow to automate daily updates to PostgreSQL and vector DB, optimizing the bot's ability to retrieve and analyze financial data for intelligent, context-driven responses.	2024
Resume-Job Matching System • Built a resume matching system using BERT embeddings to accurately rank and identify top candidate matches from resumes and job descriptions, optimizing recruitment and job search processes. • Extended the system to web-scrape job portals and public listings based on user-provided CVs and role requirements, enabling dynamic job discovery. • Implemented a multilayer, context aware RAG pipeline to combine scraped job data with resume context for improved matching relevance.	2023

PUBLICATIONS

TDISE (Text-Driven Image Synthesis and Editing), SSRN Proposed TDISE, a text-guided image synthesis and editing framework combining GAN and diffusion-based methods for precise controllable image transformations.	2024
Pixelforge: Image Revamping Through Text and Image Input Using Pix2Pix, Taylor & Francis Published a review on InstructPix2Pix-based mask-free image editing, enabling intuitive text-driven image modifications for accessible and customizable generation workflows.	2024

TRAINING AND CERTIFICATIONS

• Amazon Machine Learning Summer School '22	• IBM Qiskit Global Summer School '23
• TensorFlow Developer Professional Certificate - DL.AI	• Crash Course on Python - Google

INTERESTS

Data Engineering		Data Science		Generative AI		Quantum Computing		Particle Physics		Photography
------------------	--	--------------	--	---------------	--	-------------------	--	------------------	--	-------------